

## CNN (Convolution Neural Network) model

### 1. Model Choice and Explanation

A CNN is chosen because it can handle local patterns in steps efficiently and trains faster than many RNN variants. I consider it a good baseline before exploring RNN or LSTM.

### 2. Hyperparameters

## Keras Model

## Hyperparameters

### Initial Hyperparameters:

- Epochs = 10
- Batch size = 32
- n\_hidden (filters for Conv1D or hidden units for RNN) = 64
- Activation (hidden layers) = "relu"
- Final layer activation = "softmax"
- Optimizer = "adam"
- Loss function = "categorical\_crossentropy"

### Possible Changes:

- Might increase `n_hidden` to 128 or 256 if accuracy is too low.
- Might increase `epochs` to 20 or 30 for better training.

```
: # Basic hyperparameters
epochs = 10          # start small for testing
batch_size = 32
n_hidden = 64        # or 128, 256, etc.

timesteps = X_train.shape[1] # 15
input_dim = X_train.shape[2] # 9
n_classes = y_train.shape[1] # 15
```

```

: model = Sequential()
  # Example: CNN with kernel_size=2
  model.add(Conv1D(filters=n_hidden, kernel_size=2, activation='relu',
                   input_shape=(timesteps, input_dim)))
  model.add(MaxPooling1D(pool_size=2))

  model.add(Flatten())
  model.add(Dense(32, activation='relu'))          # a small Dense layer
  model.add(Dense(n_classes, activation='softmax'))
  # Alternative final activations if multi-label or different constraints:
  # - 'sigmoid' (common in multi-label classification)
  # - 'tanh'
  # - 'softmax' (common in multi-class classification)

  model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 14, 64)	1,216
max_pooling1d (MaxPooling1D)	(None, 7, 64)	0
flatten (Flatten)	(None, 448)	0
dense (Dense)	(None, 32)	14,368
dense_1 (Dense)	(None, 15)	495

Total params: 16,079 (62.81 KB)

Trainable params: 16,079 (62.81 KB)

Non-trainable params: 0 (0.00 B)

```
# Sklearn confusion matrix
from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test_int, y_pred_int)
print("Confusion Matrix:\n", cm)

# Station names
stations = {
    0: 'BASEL', 1: 'BELGRADE', 2: 'BUDAPEST', 3: 'DEBILT',
    4: 'DUSSELDORF', 5: 'HEATHROW', 6: 'KASSEL', 7: 'LJUBLJANA',
    8: 'MAASTRICHT', 9: 'MADRID', 10: 'MUNCHENB', 11: 'OSLO',
    12: 'SONNBLICK', 13: 'STOCKHOLM', 14: 'VALENTIA'
}

y_test_names = pd.Series([stations[i] for i in y_test_int], name='True')
y_pred_names = pd.Series([stations[i] for i in y_pred_int], name='Pred')
ctab = pd.crosstab(y_test_names, y_pred_names)
print("\nCrosstab:\n", ctab)

# 3) MEASURE FINAL ACCURACY
accuracy = accuracy_score(y_test_int, y_pred_int)
print("\nFinal Accuracy (sklearn):", accuracy)
```

180/180 ————— 1s 3ms/step

Confusion Matrix:

```
[[ 0 2599  0  0 550  0  0  0  0  0 424  0  0 109]
 [ 0 1091  0  0  1  0  0  0  0  0  0  0  0  0]
 [ 0  214  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  82  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  29  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  81  0  0  1  0  0  0  0  0  0  0  0  0]
 [ 0  11  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  61  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  9  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0 401  0  0 52  0  0  0  0  0  4  0  0  1]
 [ 0  8  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  5  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  4  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  1  0  0  0  0  0  0  0  0  0  0  0  0]]
```

Crosstab:

Pred	BELGRADE	DUSSELDORF	MUNCHENB	VALENTIA
True				
BASEL	2599	550	424	109
BELGRADE	1091	1	0	0
BUDAPEST	214	0	0	0
DEBILT	82	0	0	0
DUSSELDORF	29	0	0	0
HEATHROW	81	1	0	0
KASSEL	11	0	0	0
LJUBLJANA	61	0	0	0
MAASTRICHT	9	0	0	0
MADRID	401	52	4	1
MUNCHENB	8	0	0	0
OSLO	5	0	0	0
STOCKHOLM	4	0	0	0
VALENTIA	1	0	0	0

Final Accuracy (sklearn): 0.19013593586615546

## CNN Trials & Observations

### 1. Initial Run

- **Hyperparameters:**
  - `n_hidden = 64`, `epochs = 10`, `batch_size = 32`
  - Activation (Conv layer) = "**relu**", final layer = "**softmax**"
- **Loss:** Ranged from ~1276 at epoch 1 to ~81105 by epoch 10. (It grew quite large over epochs.)
- **Accuracy:** Reached around **19%** on the test set.
- **Stations recognized:** Roughly **4** had non-zero diagonal in the confusion matrix (BELGRADE, DUSSELDORF, MUNCHENB, VALENTIA).
- **Observation:** The model mostly predicts a few classes (huge confusion-matrix counts in those columns).

### 2. Increase `n_hidden` to 128, `epochs = 20`

- **Hyperparameters:**
  - `n_hidden = 128`, `epochs = 20`, `batch_size = 32`
  - Activation (hidden) = "**relu**", final = "**softmax**"
- **Loss:** Still grows quite large (reaches the 100k+ range, sometimes 200k+).
- **Accuracy:** ~**19%** in one run, or ~**9.5%** in another, depending on the specific run.
- **Stations recognized:** Usually sees a heavy skew toward "MADRID" or "BELGRADE," with many stations never predicted.
- **Observation:** Doubling the hidden units did not significantly improve recognition of all stations—some confusion matrices show even **lower** accuracy.

### 3. Increase `n_hidden` to 260, `epochs = 20`

- **Hyperparameters:**
  - `n_hidden = 260`, `epochs = 20`, `batch_size = 32`
  - Activation = "**relu**", final = "**softmax**"
- **Loss:** Continues to escalate into the 100k+ or 600k+ range.
- **Accuracy:** Dropped to about **7.9–9%** in the final runs.
- **Stations recognized:** The confusion matrix is dominated by one predicted class (e.g., "MADRID"), with many rows having zero correct predictions.

- **Observation:** The bigger filter count alone didn't help; the model remains heavily biased, possibly due to unscaled inputs or class imbalance.

### Conclusions / Next Steps

- **Softmax** activation is used for the final layer throughout, which is correct for multi-class classification. However, large losses suggest the network struggles.
- **Increasing filters** ( $n_{\text{hidden}}$  from 64  $\rightarrow$  128  $\rightarrow$  260) *did not* straightforwardly improve accuracy.
- **Accuracy** ranges from about **19%** at best to **7–9%** at worst. Most predictions go to a single station (e.g. BELGRADE or MADRID).
- **Potential reason:** data may need **scaling** or more **epochs**, or you might need **class weighting** for imbalanced data.
- **Stations recognized:** No run so far fully recognizes all 15 stations; often only a handful appear on the diagonal.